

IAML Notes

Lecture 1

- Supervised Learning
 - Predict output y when given input x
 - Categorical y : classification
 - Real-valued y : regression
- Unsupervised Learning
 - Internal representation of the input e.g. clustering, dimensionality
 - Getting labels is difficult and expensive
- Other areas of ML
 - Learning to predict structured objects e.g. trees, graphs
 - Reinforcement learning - learning from “rewards”
 - Semi-supervised learning - combination

Classification

No text !!!

Regression

$$f(x) = w_0 + w_1x_1 + w_Dx_D$$

$$x = (x_1, \dots, x_D)^T$$

- $f(x)$ is a linear function
- Setting of w_1, \dots, w_D is done by minimising the cost function
- Usual score function is $\sum_{i=1}^n (y^i - f(x^i))^2$.

Clustering

No text !!!

Structure of Supervised algorithms

- Define the task
- Decide on the model structure (choice of inductive bias)
- Decide on the score function (judge quality of fitted model)
- Decide on optimization/search method to optimize the score function

Supervised learning is inductive, i.e. we make generalizations about the form of $f(x)$ based on instances D .

Learning is impossible without making assumptions about f !!

Conditional Probability

$$p(X = x|Y = y) = \frac{p(x, y)}{p(y)}$$
$$p(X, Y) = p(Y)p(X|Y) = p(X)p(Y|X)$$

From the product rule,

$$p(Y|X) = \frac{p(X|Y)p(Y)}{p(X)}$$

From the sum rule the denominator is

$$p(X) = \sum_y p(X|Y)p(Y)$$

Say that Y denotes a class label, and X an observation. Then $p(Y)$ is the prior distribution for a label, and $p(Y|X)$ is the posterior distribution for Y given a data point x .

Conditional independence

$$p(X1|X2, Y) = p(X1|Y)$$

Marginal independence

$$p(X|Y) = P(X)$$

Continuous Random Variables

$$\mu = \int xp(x)dx$$
$$\sigma^2 = \int (x - \mu)^2 p(x)dx$$

For numerical discrete variables, convert integrals to sums.

Gaussian distribution

1D

$$p(x|\mu, \sigma) = N(x; \mu, \sigma) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left\{-\frac{(x - \mu)^2}{2\sigma^2}\right\}$$

- $\sqrt{2\pi\sigma^2}$ is the normalisation constant

More D

$$p(x) = \frac{1}{(2\pi)^{d/2} |\Sigma|^{1/2}} \exp\left\{-\frac{1}{2}(x - \mu)^T \Sigma^{-1} (x - \mu)\right\}$$

- Σ is the covariance matrix

$$\Sigma = E[(x - \mu)(x - \mu)^T]$$
$$\Sigma_{ij} = E[(x_i - \mu_i)(x_j - \mu_j)]$$

Likelihood

The probability of data D given model M is $p(D|M)$. This is called the **likelihood**.

$$p(D|M) = \prod_{i=1}^N p(x_i|M)$$

i.e. the product of generating each data point individually.

This is a result of the independence assumption. Try different models, pick the one with highest likelihood
-> **Maximum likelihood approach**.

Bernoulli distribution

To find most probably number of ones in n tosses, differentiate with respect to p , to find when it is maximised. Useful to take a log - it is monotonic.

$$\prod_i p(x_i|\mu, \sigma^2) = -\frac{1}{2} \sum_i \frac{(x_i - \mu)^2}{\sigma^2} - \frac{n}{2} \log(2\pi\sigma^2)$$

$$\mu = \frac{\sum_i x_i}{n}$$

$$\mu = \frac{\sum_i (x_i - \mu)^2}{n}$$

Lecture 2

Attribute-value pairs

- unordered bag of features, if structure essential, embed it
- categorical - red, yellow, blue
 - mutually exclusive
 - only equality meaningful
- ordinal - bad, good, better, best
 - there is natural ordering
 - comparison is meaningful, maths isn't
- numerical - 1, 3, 8, 4
 - usually need to normalise, remove outliers
 - maths is meaningful

Skewed distributions

- Affects regression, kNN, NB; but not DTs
- fix: take a log or atan, then normalise
- cumulative distribution function $x' = F(x) = P(X < x)$

Non-monotonic effect of attributes

- age -> probability of boxing win (not too young, not too old)
- affects regression NB, DTs (gain); less important for kNN
- if there's an obvious way to make monotonic: use it

Picking attributes - instances in the same class y should have representations x that are somehow “similar”.

Object recognition - segment image into regions, compute features describing the region.

Dealing with structure - Represent path from root to leaf as a feature.

Outliers - Isolated instances of a class that are unlike any other instances observed. Remove them or threshold.

Missing values - Try to understand why, if categorical have special category, if numerical then fill in mean or remove them. Some learners handle it.

Lecture 3

Bayesian classification

$$P(y|x) = \frac{P(x|y)P(y)}{\sum_{y'} P(x|y')P(y')}$$

- $P(y)$: prior probability of each class
 - encodes how which classes are common, which are rare
 - apriori much more likely to have common cold than Avian flu
- $P(x|y)$: class-conditional model
 - describes how likely to see observation x for class y
 - assuming it's Avian flu, do the symptoms look plausible?
- $P(x)$: normalize probabilities across observations
 - does not affect which class is most likely ($\arg \max$)

An outlier has a low probability under every class.

Independence assumption - assume $x_1, x_2 \dots$ conditionally independent given y .

$$P(x_1 \dots x_d | y) = \prod_{i=1}^d P(x_i | x_1 \dots x_{i-1}, y) = \prod_{i=1}^d P(x_i | y)$$

Probabilities of going to the beach and getting a heat stroke are not independent. May be independent if we know the weather is hot. Hot weather *explains* the dependence between beach and heatstroke. Thus, **class value explains all the dependence between attributes**.

Decision boundary

- Different means, same variance: straight line / plane
- Same mean, different variance: circle / ellipse
- General case: parabolic curve

Problems with Naive Bayes

- Zero frequency problem: add-one (Laplace) smoothing

- Word independence: add lots of hammy words to spam email
- Unable to handle correlated data - double weight to “same” features

Good stuff about Naive Bayes

- Missing data: leave out attribute that’s missing (independence assumption).
- Good complexity
 - $O(nd + cd)$ training time complexity
 - $O(dc)$ space complexity
 - $O(d)$ insertion / deletion - store partial sums instead of mean/variance
 - c - classes, n - instances, d - attributes

Naive Bayes structure

- Task: c -class classification
- Model: $c \times d$ independent assumptions
 - Continuous: Gaussian, Discrete: Bernoulli
- Score function: class-conditional likelihood
- Optimisation: analytic solution

Lecture 4

Decision trees - try to *understand* when John plays tennis. Split into sets, are they pure? If not repeat, if yes then done and see which subset new data falls into.

ID3 Algorithm

1. $A \leftarrow$ the best attribute for splitting the examples
2. Decision attribute for this node $\leftarrow A$
3. For each value of A , create new child node
4. Split training examples to child nodes
5. If examples perfectly classified then stop else iterate over new child nodes.

Which attribute to split on?

Want to measure purity of the split so we are more certain after the split. We use entropy:

$$H(S) = -p_{(+)} \log_2 p_{(+)} - p_{(-)} \log_2 p_{(-)}$$

- S : subset of training examples
- $p_{(+)}/p_{(-)}$: % of positive/negative examples in S

Information Gain

We want many items in pure sets and expect drop in entropy after split:

$$Gain(S, A) = H(S) - \sum_{V \in Values(A)} \frac{|S_V|}{|S|} H(S_V)$$

Overfitting

We can always classify training examples perfectly - keep splitting until each node has 1 example. Doesn't work on new data.

To avoid this, we stop splitting when not statistically significant. Grow the tree, then post-prune based on a validation set.

Sub-tree replacement pruning (WF 6.1)

- for each node
 - pretend to remove node + all children
 - measure performance on validation set
- remove node that results in greatest improvement
- repeat until harmful

Structure of DT

- Task: classification, discriminative
- Model: decision tree
- Score function: information gain at each node, prefer short trees and high-gain near root
- Optimisation: greedy search from simple to complex guided by information gain

Problems of DT

- Biased towards attributes with many values
- Doesn't work on new (un-observed) data
- Only axis aligned splits of data

GainRatio

Idea is to penalise attributes with many values.

$$SplitEntropy(S, A) = - \sum_{V \in Values(A)} \frac{|S_V|}{|S|} \log \frac{|S_V|}{|S|}$$
$$GainRatio(S, A) = \frac{Gain(S, A)}{SplitEntropy(S, A)}$$

Continuous DTs - Create ranges to split on which can be optimised (WF 6.1).

Multi-class classification

- predict most frequent class
- Entropy: $H(S) = - \sum_c p_{(c)} \log_2 p_{(c)}$
- $p_{(c)}$ % of examples of class c in S

Regression

- predicted output: mean of the training examples in subset
- requires a different definition of entropy (dunno)
- can use linear regression at the leaves (WF 6.5)

Good stuff about DTs

- interpretable
- easily handles irrelevant attributes (Gain = 0)
- can handle missing data - (WF 6.1)
- very compact # of nodes $\ll d$ after pruning
- very fast at testing time $O(depth)$

Random decision forest

- Grow K different decision trees
 - Pick a random subset S_r of training examples
 - grow a full ID3 tree T_r (no pruning)
 - * when splitting pick from $d \ll D$ random attributes
 - * compute gain based on S_r instead of full set
 - repeat for $r = 1 \dots K$
- Given new data point X
 - classify X using each tree and then use majority vote (state-of-the-art performance)