

Parallel Architectures

- Shared Memory
  - UMA
  - NUMA
- Multicomputer
  - like NUMA but no sharing only message passing

Consistency Models

- Sequential Consistency: every CPU sees the same interleaving
  - Program order
  - Actions of different threads interleaved arbitrarily
- Release Consistency: writes visible after sync

Parallel Patterns

- Bag of Tasks
  - One bag with synced access
  - Collection of bags with task-stealing
- Pipeline: sequential producer-consumer relationship
- Interacting Peers: info exchanged in both directions

Shared Variable Synchronisation

- Mutex: only one accesses a resource at a time
- Condition Synchronisation: delay until condition met
- Deadlock: two or more trying to enter CS, none succeeds
- Absence of Delay: if only one, thread is not prevented
- Eventual Entry: thread attempting to access CS succeeds

Locks

- Shared boolean variable with await: requires SC
- Test-and-Set: too many atomic TS operations
- Test-and-Test-and-Set: loop on cached, check again
- Bakery Algorithm: requires SC

Barriers

- Sense Reversing Barrier
- Symmetric Barrier
- Dissemination Barrier

Structured Primitives

- Semaphores
  - Buffer: binary semaphore
  - Bounded Buffer: n-semaphore
  - Producers-Consumers: n-semaphore + 2 bin. semaphores
- Monitors: one thread active within a method at any time

Pthreads

- threads
  - type: pthread\_t
  - pthread\_create(tid, attr, f, arg)
  - pthread\_join(tid, result)
- semaphores
  - type: sem\_t
  - sem\_init(sem, share, init)
  - sem\_wait(&s):  $P(s)$
  - sem\_post(&s):  $V(s)$
- Mutex
  - type: pthread\_mutex\_t
  - pthread\_mutex\_init(&m, attr)

- pthread\_mutex\_lock(&m)/pthread\_mutex\_trylock(&m)
  - pthread\_mutex\_unlock(&m)
- Condition Variables
  - type: pthread\_cond\_t
  - pthread\_cond\_wait(&cv, &mut): mut is lock already held
  - pthread\_cond\_signal(&cv)
  - pthread\_cond\_broadcast(&cv)

Java

Messaging

- Synchronisation: async/sync
- Addressing: wildcard/destination
- Collective Ops: point-to-point or not
  - Broadcast: everyone gets a copy of same value
  - Scatter: everyone gets a portion of an array
  - Gather: opposite of scatter
  - Reduction: combine gathered values at one node
  - Scan (Prefix): Reduction with partials as well

MPI Concepts

- Communicator: set of processes
- MPI\_Init(argc, argv) - before any MPI calls
- MPI\_Finalize() - after all MPI calls
- MPI\_COMM\_WORLD - global communicator
- MPI\_Comm\_size - size of communicator
- MPI\_Comm\_rank - rank within communicator
- MPI\_Comm\_spawn(command, argv, p, info, root, comm, intercomm, errcodes) - create p new processes with intercomm being MPI\_COMM\_WORLD for new processes
- MPI\_Comm\_get\_parent(comm) - gets the communicator used in spawn

MPI Sends and Receives

- MPI\_Send(buf, count, type, dest, tag, comm) - send count sized chunk of type type stored in buf to dest in communicator comm with (non-negative) id tag
- MPI\_Recv(buf, count, type, src, tag, comm, status) - receive count sized chunk of type type into buf from src in comm with id tag, store status into status
  - MPI\_ANY\_SOURCE and/or MPI\_ANY\_TAG can be used as wildcards
- A receive will match send with matching communicator, tag, and source in order
- Status info such as s.MPI\_SOURCE and s.MPI\_TAG

**Blocking:** Returns only when it is safe to reuse the buffer. **Not** after a matching operation has been executed.

Types of blocking communication

- MPI\_Rsend works only if MPI\_Recv already called
- MPI\_Ssend returns only after matching receive found
- MPI\_Send returns after buffer safe to reuse blocking if buffer full
- MPI\_Bsend same as previous, error if buffer full
- MPI\_Recv blocks until message completely in buffer

Types of non-blocking communication

- Put I in front of everything
- Extra request parameter - handle for MPI\_Wait and MPI\_Test to wait or check for completion
- MPI\_Isend returns before buffer safe for reuse

## Message Probing

- `MPI_Probe(src, tag, comm, status)` - fills status without receiving, inspect buffer size
- `MPI_Iprobe(src, tag, comm, flag, status)` - flag signals whether message available
- `MPI_Get_count(status, type, count)` - sets count to number of `type` items described by status

## Collective Operations

- `MPI_Bcast(buf, count, type, root, comm)` - broadcast of `buf` from `root` to all processes in `comm`
- `MPI_Scatter(sendbuf, sendcount, sendtype, recvbuf, recvcount, recvtype, root, comm)` - `i`-th chunk of size `sendcount` from `sendbuf` is sent to `recvbuf`
- `MPI_Allreduce(sendbuf, recvbuf, count, sendtype, op, comm)` - reduces `sendbuf` from all within `comm` point-wise into everyone's `recvbuf`

## Communicator Splitting:

: `MPI_Comm_split(old, colour, key, new)` - processors are split based on whether `key == colour`.

## TBB

- Task scheduler manages thread pool with task stealing
- `parallel_for(range, body)`
  - Range
    - \* copy constructor and a destructor
    - \* defines `is_empty`
    - \* defines `is_divisible`
    - \* defines splitting constructor `R(R &r, split)`
    - \* predefined ranges: `blocked_range`, `blocked_range2d`
  - Body
    - \* copy constructor and a destructor
    - \* defines `operator`
  - `parallel_for(0, N, addone)` - runs `addone` for `i` in `0..N`
- `parallel_reduce`
  - Range: same as `parallel_for`
  - Body
    - \* splitting constructor and destructor
    - \* defines `operator`
    - \* defines `join`

## Linda

- global, content-addressable memory
- processes run asynchronously, have six operations
- `out(exp1, exp2, ..., expN)` - add to tuple space
- `in("Green", ?y, ?r, FALSE)` - take from tuple space
- `rd(template)` - check existence of tuple
- `eval(exp1, exp2, ..., expN)` - add to tuple space, `expX` ran in separate thread
- non-blocking `inp`, `rdp` - return immediately